

L. Lancia

G. Salillari

Cloud Computing

Master Degree in Data Science

Sapienza Università di Roma

Facebook Tao

Distributed Data Store for the Social Graph

Table of Contents

Introduction

What is TAO?

Tao

is a geographically distribute store

- deployed at Facebook
- with efficient and timely access to social graph
- using a fixed set of query
- replacing memcache
- running on thousands of machines
- provide access to many PB of data
- process a billion reads ad millions of writes each second!

The social graph

Facebook has more than 1 billion active user

- recording relationships,
- sharing interests,
- uploading pictures and ...

The user experience of Fb comes from rapid, efficient and scalable access to the *social graph*

What's behind an entry in your Fb page?

A single Fb page aggregate and filter hundreds of items from the social graph.

producing user-tailored content and checking for privacy

